

# Chapter 3

Wenduo Wang

July 20, 2016

## Problem 15

- (a) Let's create a simple linear regression model for each variable against `crim`, and plot the fitted line and actual data points.

```
library(MASS)
library(dplyr)
```

```
>>>
```

```
>>> Attaching package: 'dplyr'
```

```
>>> The following object is masked from 'package:MASS':
```

```
>>>
```

```
>>>     select
```

```
>>> The following objects are masked from 'package:stats':
```

```
>>>
```

```
>>>     filter, lag
```

```
>>> The following objects are masked from 'package:base':
```

```
>>>
```

```
>>>     intersect, setdiff, setequal, union
```

```
summary(Boston)
```

```
row_no <- nrow(Boston)
```

```
training_rows <- sample(1:row_no, round(row_no*0.8))
```

```
training_set <- Boston[training_rows,]
```

```
test_set <- setdiff(Boston, training_set)
```

```
confint_df <- data.frame()
```

```
par(mfrow=c(4, 4), mar=c(2, 1, 2, 1))
```

```
for (i in c(1:ncol(Boston))) {
```

```
  if (colnames(Boston)[i]=="crim"){
```

```
    next
```

```
  }
```

```
  formula <- as.formula(paste("crim~", colnames(Boston)[i]))
```

```
  lm_model <- lm(formula, data=training_set)
```

```
  plot(y=Boston$crim, x=Boston[,i], main=paste("crim ~", colnames(Boston)[i]), pch=1, col="lightgray")
```

```
  abline(lm_model$coef[1], lm_model$coef[2], col="red")
```

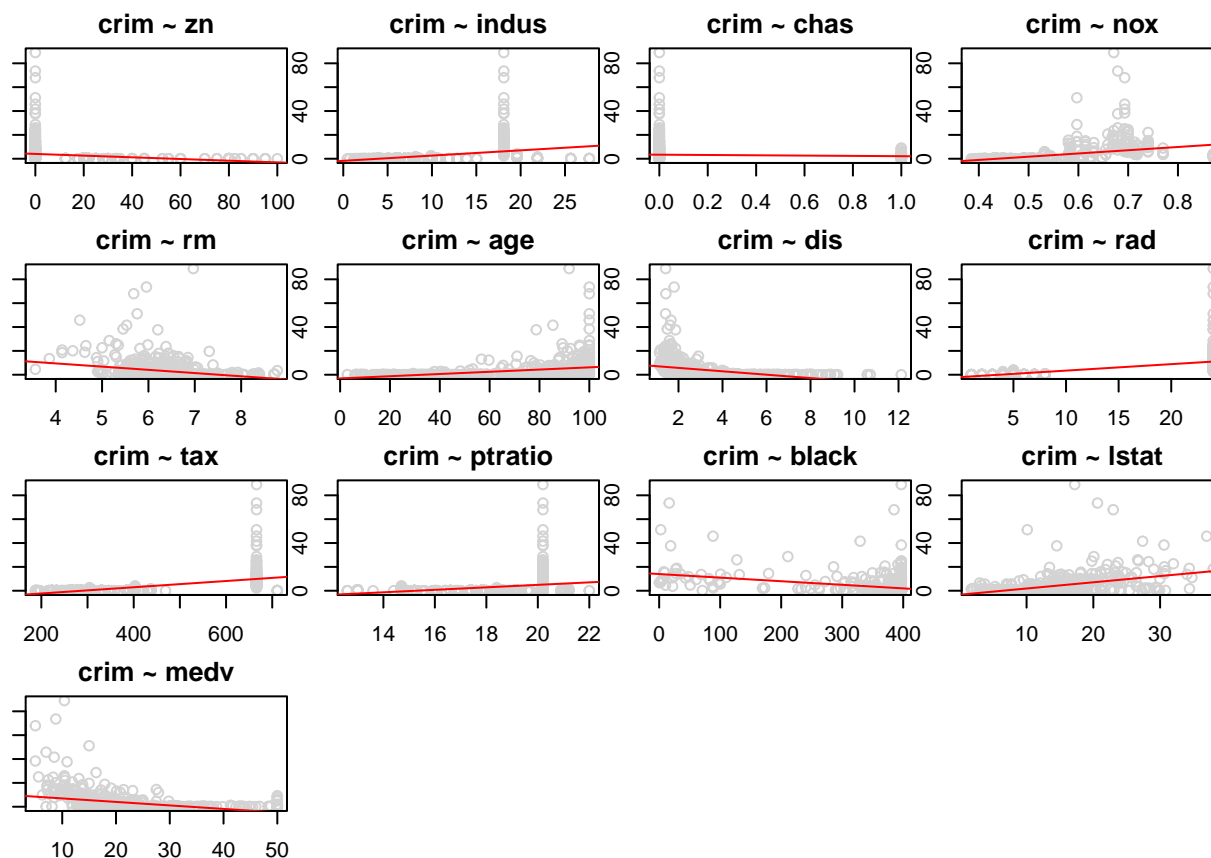
```
  confint_df <- rbind(confint_df, confint(lm_model)[2,])
```

```
}
```

```
print(dim(confint_df))
```

```
colnames(confint_df) <- c("2.5%", "97.5%")
```

```
row.names(confint_df) <- colnames(Boston)[colnames(Boston)!="crim"]
```



```
>>>      crim          zn          indus          chas
>>> Min.   : 0.00632   Min.   : 0.00   Min.   : 0.46   Min.   :0.00000
>>> 1st Qu.: 0.08204   1st Qu.: 0.00   1st Qu.: 5.19   1st Qu.:0.00000
>>> Median : 0.25651   Median : 0.00   Median : 9.69   Median :0.00000
>>> Mean    : 3.61352   Mean    : 11.36   Mean    :11.14   Mean    :0.06917
>>> 3rd Qu.: 3.67708   3rd Qu.: 12.50   3rd Qu.:18.10   3rd Qu.:0.00000
>>> Max.    :88.97620   Max.    :100.00   Max.    :27.74   Max.    :1.00000

>>>      nox          rm          age          dis
>>> Min.   :0.3850   Min.   :3.561   Min.   : 2.90   Min.   : 1.130
>>> 1st Qu.:0.4490   1st Qu.:5.886   1st Qu.: 45.02   1st Qu.: 2.100
>>> Median :0.5380   Median :6.208   Median : 77.50   Median : 3.207
>>> Mean    :0.5547   Mean    :6.285   Mean    : 68.57   Mean    : 3.795
>>> 3rd Qu.:0.6240   3rd Qu.:6.623   3rd Qu.: 94.08   3rd Qu.: 5.188
>>> Max.    :0.8710   Max.    :8.780   Max.    :100.00   Max.    :12.127

>>>      rad          tax          ptratio          black
>>> Min.   : 1.000   Min.   :187.0   Min.   :12.60   Min.   : 0.32
>>> 1st Qu.: 4.000   1st Qu.:279.0   1st Qu.:17.40   1st Qu.:375.38
>>> Median : 5.000   Median :330.0   Median :19.05   Median :391.44
>>> Mean    : 9.549   Mean    :408.2   Mean    :18.46   Mean    :356.67
>>> 3rd Qu.:24.000   3rd Qu.:666.0   3rd Qu.:20.20   3rd Qu.:396.23
>>> Max.    :24.000   Max.    :711.0   Max.    :22.00   Max.    :396.90

>>>      lstat          medv
>>> Min.   : 1.73   Min.   : 5.00
>>> 1st Qu.: 6.95   1st Qu.:17.02
>>> Median :11.36   Median :21.20
>>> Mean    :12.65   Mean    :22.53
>>> 3rd Qu.:16.95   3rd Qu.:25.00
```

```
>>> Max.      :37.97    Max.      :50.00
>>> [1] 13  2
```

The linear model does not fit the data very well according to the plots. To further assess the correlation, let's have a closer look at the coefficients of these models

```
print(confint_df)
```

```
>>>
>>>          2.5%      97.5%
>>> zn      -0.09935118 -0.04268520
>>> indus    0.35808317  0.51961541
>>> chas     -3.82768332  1.29807214
>>> nox      22.51577304 31.79801225
>>> rm       -3.55908946 -1.84409106
>>> age      0.07281828  0.11311928
>>> dis      -1.70160646 -1.14786283
>>> rad      0.49330998  0.58986045
>>> tax      0.02345440  0.02887982
>>> ptratio  0.76702904  1.31840736
>>> black    -0.03613965 -0.02384384
>>> lstat     0.45027390  0.59180300
>>> medv     -0.35537234 -0.23229956
```

The above code lists the 95% confidence intervals of coefficients for all the predictors with respect to each linear model. As seen from the result, the confident interval of the coefficient of **chas** contains zero, which means this predictor is probably not correlated with **crim**. Meanwhile, other predictors mostly have very small coefficients close to 0, with the exception of **nox**. So up to now, it appears **nox** is most likely to be a true predictor of **crim**.

(b, c) Now let's create a linear model of **crim** on all predictors, and see their coefficients' confident intervals.

```
lm_model_multiple <- lm(crim~., data=training_set)
print(confint(lm_model_multiple))
```

```
>>>
>>>          2.5 %      97.5 %
>>> (Intercept)  0.275792710 20.3372816122
>>> zn          0.005906254  0.0602465169
>>> indus       -0.185803063  0.0406832888
>>> chas        -2.181862616  1.1271352400
>>> nox        -12.797970913  1.3004008691
>>> rm          -1.181755678  0.5438208326
>>> age         -0.032856938  0.0153680742
>>> dis         -1.123715298 -0.3230844705
>>> rad          0.363724585  0.5982704137
>>> tax         -0.008207642  0.0053972468
>>> ptratio     -0.417741157  0.1023245419
>>> black       -0.009084191  0.0009468215
>>> lstat        0.142689480  0.3494199391
>>> medv        -0.118350664  0.0423924819
```

When all the predictors are simultaneously fitted against **crim**, the result turns out very different from the previous single linear regression models. First, the previously strong predictor **nox** is now largely irrelevant. Actually in this case only **zn**, **dis** and **rad** show a significant correlation with **crim**, which is equivalent to rejecting the null hypothesis that  $H_0: \beta_j=0$ .

(d) To inspect if there is a correlation in the form of

$$crim = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3$$

we create a linear regression model for each predictor against `crim`, in the format of `lm(crim ~ x + x^2 + x^3, data=training_set)`, and then inspect the model's coefficient confidence interval.

```
for (i in c(1:ncol(Boston))) {
  if (colnames(Boston)[i] == "crim") {
    next
  }
  variable <- colnames(Boston)[i]
  formula <- as.formula(paste("crim ~ ", variable, " + I(", variable, "^2) + I(", variable, "^3)", sep=""))
  cat("For predictor:", colnames(Boston)[i], "\n")
  pn_model_3 <- lm(formula, data=training_set)
  print(confint(pn_model_3))
  cat("\n")
}
```

```
>>> For predictor: zn
>>>           2.5 %      97.5 %
>>> (Intercept) 3.658861e+00 5.051352e+00
>>> zn          -4.914741e-01 -1.112041e-01
>>> I(zn^2)     -8.441882e-04 1.290484e-02
>>> I(zn^3)     -9.308368e-05 2.095429e-05
>>>
>>> For predictor: indus
>>>           2.5 %      97.5 %
>>> (Intercept) 0.631334004 5.778530428
>>> indus       -2.454347666 -0.920315743
>>> I(indus^2)  0.155025514 0.277047179
>>> I(indus^3) -0.007414645 -0.004517687
>>>
>>> For predictor: chas
>>>           2.5 %      97.5 %
>>> (Intercept) 2.746257 4.044981
>>> chas        -3.827683 1.298072
>>> I(chas^2)    NA      NA
>>> I(chas^3)    NA      NA
>>>
>>> For predictor: nox
>>>           2.5 %      97.5 %
>>> (Intercept) 160.8801 261.0944
>>> nox         -1405.4935 -899.6081
>>> I(nox^2)     1602.2452 2429.8508
>>> I(nox^3)     -1331.4259 -892.1624
>>>
>>> For predictor: rm
>>>           2.5 %      97.5 %
>>> (Intercept) 219.783988 479.3125698
>>> rm          -209.263732 -86.4206344
>>> I(rm^2)      11.328428 30.4951872
>>> I(rm^3)      -1.475121 -0.4907945
>>>
>>> For predictor: age
>>>           2.5 %      97.5 %
>>> (Intercept) -5.796322e+00 2.561622e+00
>>> age         -1.057828e-01 4.693557e-01
>>> I(age^2)     -1.055695e-02 7.467448e-04
```

```

>>> I(age^3)      8.131208e-06 7.384325e-05
>>>
>>> For predictor: dis
>>>           2.5 %      97.5 %
>>> (Intercept) 21.1065970 28.5750522
>>> dis         -15.2181882 -9.9027494
>>> I(dis^2)    1.4101874 2.4722936
>>> I(dis^3)    -0.1235113 -0.0605432
>>>
>>> For predictor: rad
>>>           2.5 %      97.5 %
>>> (Intercept) -3.402618045 2.171492473
>>> rad         -0.904057033 1.945109126
>>> I(rad^2)    -0.277490344 0.128332516
>>> I(rad^3)    -0.003174294 0.009294897
>>>
>>> For predictor: tax
>>>           2.5 %      97.5 %
>>> (Intercept) 4.962689e+00 3.887594e+01
>>> tax         -3.153024e-01 -4.145752e-02
>>> I(tax^2)    8.760718e-05 7.792764e-04
>>> I(tax^3)    -5.558744e-07 -1.872422e-08
>>>
>>> For predictor: ptratio
>>>           2.5 %      97.5 %
>>> (Intercept) 62.4440773 585.82089808
>>> ptratio     -100.9841111 -8.49334118
>>> I(ptratio^2) 0.3043651 5.70013095
>>> I(ptratio^3) -0.1050957 -0.00112827
>>>
>>> For predictor: black
>>>           2.5 %      97.5 %
>>> (Intercept) 7.433054e+00 1.465026e+01
>>> black       -3.446801e-02 1.453479e-01
>>> I(black^2)  -8.729851e-04 7.868724e-05
>>> I(black^3)  -1.886857e-07 1.198319e-06
>>>
>>> For predictor: lstat
>>>           2.5 %      97.5 %
>>> (Intercept) -2.368286644 3.524044980
>>> lstat       -0.851193726 0.485856214
>>> I(lstat^2)  -0.015414196 0.070686286
>>> I(lstat^3)  -0.001003005 0.000605381
>>>
>>> For predictor: medv
>>>           2.5 %      97.5 %
>>> (Intercept) 38.027613110 48.4434072905
>>> medv       -4.702793923 -3.3587273182
>>> I(medv^2)   0.093422906 0.1465412678
>>> I(medv^3)  -0.001433273 -0.0008062374

```

As seen from the output confidence intervals, it is first noticed that  $\text{chas}^2$  and  $\text{chas}^3$  do not have coefficients. That is due to that  $\text{chas}$  is a binary variable, whose square or cube is essentially itself, so in this case  $\text{chas}$ ,  $\text{chas}^2$  and  $\text{chas}^3$  are linearly related, and therefore the latter two polynomial terms are not fitted

in the model. `chas` put aside, several polynomial terms exhibit correlation with `crim`, whose 95% coefficient confidence intervals exclude zero. For example, `nox2` and `nox3`.