

STA 380 Homework 2

Summer Wang, Yawen Ye, Wenduo Wang

August 15, 2016

Flights at ABIA

In this dataset, we want to answer several questions: 1: Which carrier has the worst punctuality in departing time and arriving time? Which carrier has the best punctuality in departing time and arriving time? 2: Which carrier has done the best job to minimize CarrierDelay? 3: What is the best time of day to fly to minimize delays? 4: What is the best time of year to fly to minimize delays? 5: How do patterns of flights to different destinations or parts of the country change over the course of the year?


We will talk about each problem one by one.

1: Which carrier has the worst punctuality in departing time and arriving time?

Which carrier has the best punctuality in departing time and arriving time?


```
ABIA <- read.csv("data/ABIA.csv", header=T, na.strings=c("", "NA"))
attach(ABIA)
library(dplyr)
library(plotly)
```

```
p1 <- plot_ly(ggplot2::diamonds, y = ArrDelay, color = UniqueCarrier, type = "box")
p1
```




YV XE WN US UA OO OH NW MQ F9 EV DL CO B6 AA 9E


```
temp<- ABIA[which(ArrDelay!=649 & ArrDelay!= 948),]
p2 <- plot_ly(ggplot2::diamonds, y = temp$ArrDelay, color = temp$UniqueCarrier, type =
"box")
p2
```



We could actually see that OH flight is the worst when it comes to punctuality on arriving.

```
temp1<- ABIA[which(DepDelay!=665 & DepDelay!= 875),]
p3 <- plot_ly(ggplot2::diamonds, y = temp1$DepDelay, color = temp1$UniqueCarrier, type =
"box")
p3
```





We could actually see that OH and EV flight are the worst when it comes to punctuality on departing. US airline did a much better job on departure punctuality


2: Which carrier has done the best job to minimize CarrierDelay?

```
temp2<-ABIA[!is.na(CarrierDelay),]
attach(temp2)
```

```
#we could see that ArrDelay consist of five parts:
sum(temp2[, 'ArrDelay']!=temp2[, 'WeatherDelay']+temp2[, 'NASDelay']+temp2[, 'SecurityDelay']+
temp2[, 'LateAircraftDelay']+temp2[, 'CarrierDelay'])
```

```
## [1] 0
```

```
#let's get rid of the WeatherDelay, which people can not control
temp2[, 'ArrDelay_control']=temp2[, 'NASDelay']+temp2[, 'SecurityDelay']+
temp2[, 'LateAircraftDelay']+temp2[, 'CarrierDelay']
p4 <- plot_ly(ggplot2::diamonds, y = temp2[which(CarrierDelay!=875),]$CarrierDelay,
color = temp2[which(CarrierDelay!=875),]$UniqueCarrier, type = "box")
p4
```





We could see that airline MQ did the best job in term of its own control of arriving time. Airline YV did a worst job in term of its own control of arriving time. But this is actually part of the total list. So maybe, airline didn't record many records. So let's take a look at the records of each airline.

```
temp2_freq=data.frame(table(temp2$UniqueCarrier))
ABIA_freq=data.frame(table(ABIA$UniqueCarrier))
record_freq <- cbind(temp2_freq[1],temp2_freq[-1]/ABIA_freq[-1])
record_freq
```

```
##      Var1      Freq
## 1      9E 0.1647705
## 2      AA 0.2099025
## 3      B6 0.2280117
## 4      CO 0.2087757
## 5      DL 0.2708529
## 6      EV 0.2654545
## 7      F9 0.1646341
## 8      MQ 0.1794968
## 9      NW 0.2396694
## 10     OH 0.3382451
## 11     OO 0.1967621
## 12     UA 0.2304394
## 13     US 0.1207133
## 14     WN 0.1731563
## 15     XE 0.1933738
## 16     YV 0.2234682
```

```
#record_freq
p <- plot_ly(
  x = record_freq$Var1,
  y = record_freq$Freq,
  name = "record_freq",
  type = "bar")
p
```





We could see that actually MQ and YV give out the relatively good records on the five delay components. So we could say that, airline MQ did a good job on control its CarrierDelay time. YV did the worst. So for airline YV, it is important to improve its own control of arriving time


3: What is the best time of day to fly to minimize delays?

```
library(lubridate)
```

```
##
## Attaching package: 'lubridate'
```

```
## The following object is masked from 'package:base':
##
##     date
```


```
ABIA$DepTime<-format(strptime(sprintf("%04d", ABIA$DepTime), format="%H%M"), format =
"%H:%M")
ABIA$DepTime<-hm(ABIA$DepTime)
ABIA$DepTime_hour<-hour(ABIA$DepTime)
p5 <- plot_ly(ggplot2::diamonds, y = ABIA$DepDelay, color =as.character(ABIA$DepTime_hour),
type = "box")
p5
```



```

ABIA$ArrTime<-format(strptime(sprintf("%04d", ABIA$ArrTime), format="%H%M"), format =
"%H:%M")
ABIA$ArrTime<-hm(ABIA$ArrTime)
ABIA$ArrTime_hour<-hour(ABIA$ArrTime)
p6 <- plot_ly(ggplot2::diamonds, y = ABIA$ArrDelay, color =as.character(ABIA$ArrTime_ho
ur), type = "box")
p6

```




We can see from the plot that, 5:00am is the best time of the day to fly to minimize delays. In the early morning, 5:00am to 7:00am is a good period of time to fly to minimize delays.


4: What is the best time of year to fly to minimize delays?

```


p7 <- plot_ly(ggplot2::diamonds, y = ABIA$DepDelay, color =as.character(ABIA$Month), ty
pe = "box")
p7

```





```
p8 <- plot_ly(ggplot2::diamonds, y = ABIA$ArrDelay, color = as.character(ABIA$Month), type = "box")
p8
```



From the plots we could see that, September is the best month of the year to fly to minimize delays.

5: How do patterns of flights to different destinations or parts of the country change over the course of the year?

```
library(plyr)
```

```
group<- ABIA %>%
```

```
  dplyr::group_by(Month, Dest) %>%
```

```
  dplyr::summarise(length(Dest))
```





```
par(mfrow=c(2,2))
```









```
colors <- rainbow(length(unique(group$Dest)))
```









```
linetype <- c(1:length(unique(group$Dest)))
```









```
plotchar <- seq(18,18+length(unique(group$Dest)),1)
```









```
for (i in (1:length(unique(group$Dest)))) {
  temp3 <- subset(group, Dest==unique(group$Dest)[i])
  xrange <- range(1:12)
  yrange<- range(temp3$`length(Dest)` )
  plot(xrange, yrange, type="n", xlab="Month", ylab=paste("Number of Flights to ", as.ch
aracter(unique(group$Dest)[i]), sep = ""))
  lines(temp3$Month, temp3$`length(Dest)` , type="b", lwd=1.5)
  line <- readline()
}
```


































There is clearly a drop in the frequency of the flights to some destinations starting around September. These destinations include ABQ, AUS, CLE, CVG, IAD, LAS, MEM, MSY, ONT, RDU, SLC, SNA, TUS, SEA. The frequency of the flights to some destinations even drop around June, which include BNA, DAL, MAF, MCI.

Author attribution

Import `tm` and `dplyr` libraries. The former is the text mining library, and the latter provides tools for matrix operations.

```
library(compiler)
enableJIT(3)
```

```
## [1] 0
```

```
rm(list=ls())
gc()
```

```
##          used (Mb) gc trigger (Mb) max used (Mb)
## Ncells  655161 35.0   1168576 62.5   1168576 62.5
## Vcells 3284682 25.1   10405454 79.4  13004239 99.3
```

```
library(tm)
library(dplyr)
```

The first thing to do with the documents is to convert them to `Corpus` objects defined in `tm` library. While the texts are read in, a series of operations are performed, including enforcing lowercase, removing numbers & punctuations, and stripping white spaces. It is sometimes problematic to remove *stopwords* as such could change the meaning of the text. But given the abundance of data in this case, removing *stopwords* provides more benefits in terms of simplifying computation.

The training set and test set are converted into `Corpus` objects in the above way, and then the object is transformed into a `DocumentTermMatrix` and later normal matrix. No sparse terms are teased out at this stage.

First let's define a set of helper functions.


```

# a function to read articles inside a folder
readArticle <- function(article_url){
  article_list <- lapply(article_url,
    function(article_url) readPlain(
      elem=list(content=readLines(article_url)),
      language="en",
      id=article_url
    ))
  return(article_list)
}

# convert an article into Corpus
get_corpus <- function(article){
  corp <- Corpus(VectorSource(article))
  return(corp)
}

# clean up the corpus object, i.e. standardize corpus format
clean_corpus <- function(corp, tfidf=0){
  corp <- tm_map(corp, content_transformer(tolower))
  corp <- tm_map(corp, content_transformer(removeNumbers))
  corp <- tm_map(corp, content_transformer(removePunctuation))
  corp <- tm_map(corp, content_transformer(stripWhitespace))
  corp <- tm_map(corp, content_transformer(removeWords), stopwords("en"))

  # define an optional "control" parameter
  # to weight the terms according to their tfidf index
  if (tfidf) {
    control <- list(weighting=function(x) weightTfIdf(x, normalize=FALSE))
    corp_DTM <- DocumentTermMatrix(corp, control=control)
  } else {
    corp_DTM <- DocumentTermMatrix(corp)
  }

  corp_mat <- as.matrix(corp_DTM)

  return(corp_mat)
}

# a function to convert all the folders and articles inside the folders
# into a list of corpus.
# every element in the corpus list is a terms vector of a document
# inside the same folder.
get_mat <- function(url){
  author_list <- Sys.glob(url)
  folder_list <- Sys.glob(paste(author_list, "/*", sep=""))
  articles <- lapply(folder_list, readArticle)
  corpus_list <- lapply(articles, get_corpus)
  corpus_list <- lapply(corpus_list, clean_corpus)
  return(corpus_list)
}

# a function to calculate the angle between two given vecotrs
get_cosine <- function(v1, v2){
  # first measure the original length of each vector
  len1 <- sqrt(v1%*%t(v1))
  len2 <- sqrt(v2%*%t(v2))
  # only the common columns between the two matrices are kept.

```

```

v1_conform <- v1[colnames(v1) %in% colnames(v2)]
v2_conform <- v2[colnames(v2) %in% colnames(v1)]
# if the two vectors are orthogonal, then return 0
if (length(v1_conform)*length(v2_conform)==0){
  return(0)
} else {
  # vector angle formula
  cosine <- t(v1_conform) %*% v2_conform / (len1*len2)
  return(cosine)
}
}

# translate the prediction matrix to author indices
pred_author <- function(pred){
  article_index <- lapply(pred, which.max)
  author_index <- lapply(article_index, function(ind) ceiling(ind/50))
  return(author_index)
}

# calculate prediction accuracy
pred_accuracy <- function(author_list){
  target <- rep(1:50, each=50)
  len <- length(author_list)
  accuracy <- mean(author_list==target[1:len])
  return(accuracy)
}

# get term frequencies within each article for Naive Bayes
get_frequency <- function(mat){
  frequency_mat <- log(mat) / as.vector(rowSums(mat))
}

# calculate the cumulative term probability for each article
# only the common columns between the two matrices are kept
get_product <- function(v1, v2){
  v1_conform <- v1[colnames(v1) %in% colnames(v2)]
  v2_conform <- v2[colnames(v2) %in% colnames(v1)]
  if (length(v1_conform)*length(v2_conform)==0){
    return(0)
  } else {
    product <- t(v1_conform) %*% v2_conform
    return(product)
  }
}

```

Read in training and test data.

Due to high resource consumption, the following chunks are disabled.

```

train_mat <- get_mat("data/ReutersC50/C50train/*")

test_mat <- get_mat("data/ReutersC50/C50test/*")

```

Model 1: Vector angle

The first model is based on the angle between the train article terms vector and the test article terms vector. In this model, each new article from the test set is compared with all articles in the train set, by calculating the cosine of their vector angle. The train article showing the smallest angle is taken as the output, and therefore the test article is deemed to belong to the same author as the former.

```
cosine_mat <- lapply(test_mat[1:100],  
                    function(v) lapply(train_mat, function(x) get_cosine(x, v)))  
  
author_list_1 <- pred_author(cosine_mat)  
accuracy_1 <- pred_accuracy(author_list_1)  
accuracy_1
```

The accuracy is ~54%.

Let's have a look at the authors whose articles are most difficult to guess.

```
authors <- Sys.glob("data/ReutersC50/C50train/*")  
authors <- gsub("."+ "/", "", authors)  
author_list_1 <- as.matrix((read.csv("author_list_vec.csv"))[-1]))  
author_list_1 <- data.frame(t(author_list_1))  
author_df <- data.frame(author=rep(authors, each=50))  
author_df <- cbind(author_df, author_list_1[,1], rep(1:50, each=50))  
colnames(author_df) <- c("author", "prediction", "actual")  
author_df <- cbind(author_df, (author_df$prediction==author_df$actual)*1)  
colnames(author_df)[ncol(author_df)] <- "mistake"
```

```
author_agg <- aggregate(mistake~author, data=author_df, FUN="sum")  
author_agg[order(author_agg$mistake, decreasing = TRUE),]
```

```

> author mistake
> 29 LynnleyBrowning 48
> 11 FumikoFujisaki 47
> 16 JimGilchrist 47
> 21 KarlPenhaul 45
> 33 MatthewBunce 43
> 6 BradDorfman 39
> 17 JoeOrtiz 39
> 12 GrahamEarnshaw 38
> 28 LynneO'Donnell 38
> 47 TheresePoletti 37
> 1 AaronPressman 36
> 34 MichaelConnor 36
> 36 NickLouth 35
> 22 KeithWeir 34
> 14 JaneMacartney 32
> 26 KouroshKarimkhany 32
> 27 LydiaZajc 32
> 48 TimFarrand 32
> 32 MartinWolk 30
> 40 RobinSidel 30
> 41 RogerFillion 30
> 45 SimonCowell 30
> 24 KevinMorrison 29
> 37 PatriciaCommins 29
> 39 PierreTran 29
> 5 BernardHickey 28
> 19 JonathanBirt 28
> 46 TanEeLyn 27
> 38 PeterHumphrey 24
> 23 KevinDrawbaugh 23
> 31 MarkBendeich 23
> 43 SarahDavison 23
> 18 JohnMastrini 22
> 30 MarcelMichelson 22
> 49 ToddNissen 22
> 25 KirstinRidley 21
> 2 AlanCrosby 20
> 13 HeatherScofield 20
> 8 DavidLawder 19
> 10 EricAuchard 19
> 9 EdnaFernandes 18
> 20 JoWinterbottom 18
> 44 ScottHillis 16
> 50 WilliamKazer 16
> 3 AlexanderSmith 15
> 35 MureDickie 15
> 7 DarrenSchuettler 11
> 42 SamuelPerry 9
> 4 BenjaminKangLim 8
> 15 JanLopatka 6

```

Model 2: Naive Bayes

We can also apply Naive Bayes method to determine the authors. Specifically, in this case our problem is to predict the probability

$$P(\text{the author is } x | \text{given the terms matrix of an author})$$

Where x can be any of the known authors from the training set. According to the general Bayes theorem, the above probability is equal to

$$\frac{P(\text{the author's new document terms matrix is like the test set} | \text{the author is } x) \cdot P(\text{author is } x)}{P(\text{a new document terms matrix is like the test set})}$$

Since $\frac{P(\text{author is } x)}{P(\text{a new document terms matrix is like the test set})}$ is the same for all articles, we consider it as a constant, and therefore we can only calculate and compare

$$P(\text{the author's new document terms matrix is like the test set} | \text{the author is } x).$$

Take logarithm of the expression and we will get the log probability terms as matrix expression

$$\text{test} * \log \frac{\text{train}}{\text{sum}(\text{train})}$$

```
# get term frequencies for each terms vector in the training set
train_freq <- lapply(train_mat, get_frequency)
# multiply the test set article vector by that in the training set
prob_mat <- lapply(test_mat[1:100],
  function(v) lapply(train_freq, function(x) get_product(x, v)))
# select the author with the highest conditional probability for each article
author_list_2 <- pred_author(prob_mat)
accuracy_2 <- pred_accuracy(author_list_2)
accuracy_2
author_agg <- aggregate(mistake~author, data=author_df, FUN="sum")
author_agg[order(author_agg$mistake, decreasing = TRUE),]
```

The accuracy is ~47%.

Let's see the authors that were guesses wrong most.

```
author_list_2 <- as.matrix((read.csv("author_list_NB.csv"))[-1]))
author_list_2 <- data.frame(t(author_list_2))
author_df <- data.frame(author=rep(authors, each=50))
author_df <- cbind(author_df, author_list_2[,1], rep(1:50, each=50))
colnames(author_df) <- c("author", "prediction", "actual")
author_df <- cbind(author_df, (author_df$prediction==author_df$actual)*1)
colnames(author_df)[ncol(author_df)] <- "mistake"
```

```
author_agg <- aggregate(mistake~author, data=author_df, FUN="sum")
author_agg[order(author_agg$mistake, decreasing = TRUE),]
```

```

>          author mistake
> 16      JimGilchrist      49
> 21      KarlPenhaul      43
> 29      LynnleyBrowning   43
> 33      MatthewBunce     42
> 11      FumikoFujisaki   40
> 1       AaronPressman    36
> 26      KouroshKarimkhany 36
> 12      GrahamEarnshaw   34
> 17      JoeOrtiz          34
> 27      LydiaZajc        34
> 14      JaneMacartney    33
> 28      LynneO'Donnell    32
> 34      MichaelConnor    32
> 47      TheresePoletti    32
> 5       BernardHickey    30
> 9       EdnaFernandes    29
> 40      RobinSidel       29
> 41      RogerFillion     29
> 6       BradDorfman      28
> 45      SimonCowell      28
> 19      JonathanBirt     27
> 30      MarcelMichelson  27
> 39      PierreTran       27
> 37      PatriciaCommins  25
> 22      KeithWeir       24
> 23      KevinDrawbaugh   23
> 24      KevinMorrison    23
> 46      TanEeLyn        23
> 48      TimFarrand       21
> 32      MartinWolk       20
> 13      HeatherScofield  19
> 18      JohnMastrini     19
> 31      MarkBendeich     19
> 36      NickLouth        19
> 38      PeterHumphrey    19
> 20      JoWinterbottom   15
> 44      ScottHillis      15
> 49      ToddNissen       14
> 50      WilliamKazer     14
> 2       AlanCrosby       12
> 10      EricAuchard      12
> 25      KirstinRidley    12
> 8       DavidLawder      11
> 35      MureDickie       11
> 7       DarrenSchuettler  9
> 42      SamuelPerry      9
> 4       BenjaminKangLim  6
> 3       AlexanderSmith   3
> 15      JanLopatka       3
> 43      SarahDavison     3

```

It is seen that the two models give similar prediction accuracies, and the vanilla vector angle method is even a little bit more accurate than Naive Bayes, possibly suggesting dependency between terms (which is likely true).

Let's see which authors' articles are hard to predict.

Practice with association rule mining

```
library(compiler)
enableJIT(3)
```

```
## [1] 3
```

```
rm(list=ls())
gc()
```

```
##          used (Mb) gc trigger (Mb) max used (Mb)
## Ncells  679899 36.4   1168576 62.5   1168576 62.5
## Vcells 3331504 25.5   10405454 79.4  13004239 99.3
```

Import `arules` library for association rule mining.

```
library(arules)
```

Import data with `read.transactions()` function from `arules`, which will automatically convert each row into a list of items separated by commas, and the returned object is a `transactions` object. This function will also drop duplicate items from each basket if `rm.duplicates = TRUE`.

```
groceries <- read.transactions("data/groceries.txt", sep=",", rm.duplicates = TRUE)
```

We can assign each user an id by converting the `transactions` to a list with `as(from="transactions", to="list")`, and define `names()` of the list, and then convert the list back to `transactions`.

```
groceries_list <- as(groceries, "list")
names(groceries_list) <- as.character(1:length(groceries_list))
groceries <- as(groceries_list, "transactions")
```

At this point, the `groceries` object is suitable for a priori analysis. Before applying the `apriori` function, we need to determine what the *support* and *confidence* thresholds, and *maxlen* value. This is essentially a heuristic process, so here let's first try a higher *support* level 0.01 and *confidence* threshold 0.55 (just a little bit more than 0.5), and see what we get.

```
params <- list(support=.01, confidence=.55, maxlen=4)
grocery_rules <- apriori(groceries, parameter = params)
```

```

> Apriori
>
> Parameter specification:
> confidence minval smax arem aval originalSupport support minlen maxlen
>      0.55      0.1      1 none FALSE          TRUE      0.01      1      4
> target      ext
> rules FALSE
>
> Algorithmic control:
> filter tree heap memopt load sort verbose
>    0.1 TRUE TRUE  FALSE TRUE     2     TRUE
>
> Absolute minimum support count: 98
>
> set item appearances ...[0 item(s)] done [0.00s].
> set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
> sorting and recoding items ... [88 item(s)] done [0.00s].
> creating transaction tree ... done [0.00s].
> checking subsets of size 1 2 3 4 done [0.00s].
> writing ... [7 rule(s)] done [0.00s].
> creating S4 object ... done [0.00s].

```

```
inspect(subset(grocery_rules, subset=lift>=2))
```

```


> lhs                rhs                support confidence    lift
> 1 {curd,           => {whole milk}      0.01006609  0.5823529  2.279125
>   yogurt}
> 2 {butter,         => {whole milk}      0.01148958  0.5736041  2.244885
>   other vegetables}
> 3 {domestic eggs,  => {whole milk}      0.01230300  0.5525114  2.162336
>   other vegetables}
> 4 {citrus fruit,   => {other vegetables} 0.01037112  0.5862069  3.029608
>   root vegetables}
> 5 {root vegetables,=> {other vegetables} 0.01230300  0.5845411  3.020999
>   tropical fruit}
> 6 {root vegetables,=> {whole milk}      0.01199797  0.5700483  2.230969
>   tropical fruit}
> 7 {root vegetables,=> {whole milk}      0.01453991  0.5629921  2.203354
>   yogurt}

```

Here we only selected the associations with *lift* greater than 2, which gives 7 in total. And among them are items such as *other vegetables* and *whole milk*, which are themselves frequent terms across all baskets. Such results provide limited information, so we need to look closer into more interesting and less ubiquitous items.

So a natural question to be asked here is, which are the most frequent items? Let's make a plot to show the top 10 frequent terms.

```
itemFrequencyPlot(groceries, topN=10)
```

So here we can see clearly that *whole milk* and *other vegetables* are indeed frequent terms containing relatively less information.

Therefore, let's lower *support* level to 0.001 to include less often items. Also, when printing out the associations, we raise the *lift* threshold to 10, which indicates highly correlated and dependent items.

```
params <- list(support=.001, confidence=.55, maxlen=4)
grocery_rules <- apriori(groceries, parameter = params)
```

```

> Apriori
>
> Parameter specification:
> confidence minval smax arem aval originalSupport support minlen maxlen
>      0.55      0.1      1 none FALSE          TRUE    0.001      1      4
> target      ext
> rules FALSE
>
> Algorithmic control:
> filter tree heap memopt load sort verbose
>    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
>
> Absolute minimum support count: 9
>
> set item appearances ...[0 item(s)] done [0.00s].
> set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
> sorting and recoding items ... [157 item(s)] done [0.00s].
> creating transaction tree ... done [0.00s].
> checking subsets of size 1 2 3 4 done [0.01s].
> writing ... [3314 rule(s)] done [0.00s].
> creating S4 object ... done [0.00s].

```

```
inspect(subset(grocery_rules, subset=lift>=10))
```

	lhs	rhs	support	confidence	lift
> 1	{liquor,				
>	red/blush wine}	=> {bottled beer}	0.001931876	0.9047619	11.23527
> 2	{popcorn,				
>	soda}	=> {salty snack}	0.001220132	0.6315789	16.69779
> 3	{Instant food products,				
>	soda}	=> {hamburger meat}	0.001220132	0.6315789	18.99565
> 4	{ham,				
>	processed cheese}	=> {white bread}	0.001931876	0.6333333	15.04549
> 5	{baking powder,				
>	flour}	=> {sugar}	0.001016777	0.5555556	16.40807
> 6	{hard cheese,				
>	whipped/sour cream,				
>	yogurt}	=> {butter}	0.001016777	0.5882353	10.61522
> 7	{hamburger meat,				
>	whipped/sour cream,				
>	yogurt}	=> {butter}	0.001016777	0.6250000	11.27867

Here we see some intriguing associations which are less frequent among all baskets but exhibits huge correlation in terms of *lift*, which is a measure of dependence. While in the previous case there were only 7 associations even with *lift* level higher than 2, here there are 7 associations with *lift* higher than 10.

Let's look at the association {liquor, red/blush wine} => {bottled beer}, it is intuitively this is some combination appealing to an alcohol lover. This intuition also holds for other association groups, such as {baking powder, flour} => {sugar} which is probably a part of common baking recipe.

And what about other associations?

```

params <- list(support=.001, confidence=.55, maxlen=4)
grocery_rules <- apriori(groceries, parameter = params)

```

```
> Apriori
>
> Parameter specification:
> confidence minval smax arem aval originalSupport support minlen maxlen
>      0.55      0.1      1 none FALSE          TRUE   0.001      1      4
> target      ext
> rules FALSE
>
> Algorithmic control:
> filter tree heap memopt load sort verbose
>    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
>
> Absolute minimum support count: 9
>
> set item appearances ...[0 item(s)] done [0.00s].
> set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
> sorting and recoding items ... [157 item(s)] done [0.00s].
> creating transaction tree ... done [0.00s].
> checking subsets of size 1 2 3 4 done [0.01s].
> writing ... [3314 rule(s)] done [0.00s].
> creating S4 object ... done [0.00s].
```

```
inspect(subset(grocery_rules, subset=(lift<=10 & lift>8)))
```

	lhs	rhs	support	confidence	lift
> 1	{frozen vegetables, specialty chocolate}	=> {fruit/vegetable juice}	0.001016777	0.6250000	8.645394
> 2	{frozen fish, other vegetables, tropical fruit}	=> {pip fruit}	0.001016777	0.6666667	8.812724
> 3	{flour, root vegetables, whole milk}	=> {whipped/sour cream}	0.001728521	0.5862069	8.177794
> 4	{misc. beverages, other vegetables, tropical fruit}	=> {fruit/vegetable juice}	0.001016777	0.5882353	8.136841
> 5	{citrus fruit, fruit/vegetable juice, grapes}	=> {tropical fruit}	0.001118454	0.8461538	8.063879
> 6	{fruit/vegetable juice, grapes, tropical fruit}	=> {citrus fruit}	0.001118454	0.6875000	8.306588
> 7	{citrus fruit, grapes, tropical fruit}	=> {fruit/vegetable juice}	0.001118454	0.6111111	8.453274
> 8	{butter, hard cheese, yogurt}	=> {whipped/sour cream}	0.001016777	0.6250000	8.718972
> 9	{butter, hard cheese, other vegetables}	=> {whipped/sour cream}	0.001220132	0.6000000	8.370213
> 10	{butter, hard cheese, whole milk}	=> {whipped/sour cream}	0.001423488	0.6666667	9.300236
> 11	{ham, other vegetables, tropical fruit}	=> {pip fruit}	0.001626843	0.6153846	8.134822
> 12	{butter, sliced cheese, whole milk}	=> {whipped/sour cream}	0.001220132	0.6000000	8.370213
> 13	{cream cheese, sugar, whole milk}	=> {domestic eggs}	0.001118454	0.5500000	8.668670
> 14	{curd, sugar, yogurt}	=> {whipped/sour cream}	0.001016777	0.6250000	8.718972
> 15	{butter, other vegetables, sugar}	=> {whipped/sour cream}	0.001016777	0.7142857	9.964539
> 16	{citrus fruit, cream cheese, whole milk}	=> {domestic eggs}	0.001626843	0.5714286	9.006410
> 17	{domestic eggs, frankfurter, tropical fruit}	=> {pip fruit}	0.001016777	0.6250000	8.261929
> 18	{shopping bags, tropical fruit, whipped/sour cream}	=> {pip fruit}	0.001118454	0.6470588	8.553526

Above are associations with *lift* between 8 and 10. And here we can see some interesting combinations such as {butter, hard cheese, milk} => {whipped/sour cream}. Why is the customer buying such protein and fat heavy foods altogether? Probably it is simply because of the way these products are placed in the store. If some products are

placed together, then they are more likely to be sold in a bundle. This can also be seen in {citrus fruit, grapes, tropical fruit} => {fruit/vegetable juice}.